A Case Study Report on

"Online Retail Transaction Clustering Using Machine Learning Techniques
" and
"Predicting Student Dropout and Academic Success Using Data Science"

for the course

Fundamentals Of Data Science (22CSC59)

Submitted in partial fulfillment of the requirements for the award of degree

BACHELOR OF ENGINEERING

in

COMPUTER SCIENCE AND ENGINEERING

by

Bhukya Vaishali (160124733142) CSE-III , 3rd Year

Faculty:

Smt. S.Durga Devi Assistant Professor



Department of Computer Science and Engineering,

Chaitanya Bharathi Institute of Technology (Autonomous),

(Affiliated to Osmania University, Hyderabad) Hyderabad, TELANGANA (INDIA) –500 075 [2026-2027]

Online Retail Transaction Clustering Using Machine Learning Techniques

1. PROBLEM STATEMENT

The online retail store maintains transaction records containing products purchased by customers. However, the store lacks a systematic approach to understanding transaction characteristics and identifying groups of similar shopping baskets. This case study aims to analyze retail transaction data using descriptive statistics, data visualization, dimensionality reduction, and clustering algorithms. K-means, hierarchical clustering, DBSCAN, and spectral clustering will be implemented and evaluated using appropriate performance metrics. The findings will help identify transaction patterns and provide insights that may support product recommendations, promotional strategies, and inventory planning.

2. METHODOLOGY

1. Load the transaction file and inspect its dimensions and schema.
2. Treat each row as one transaction and each non-empty cell as a product.
3. Normalize product text and remove repeated items within a basket for binary presence/absence analysis.
4. Compute transaction-size and product-frequency statistics.
5. Convert baskets to a binary transaction–product matrix.
6. Use PCA for a two-dimensional visualization.
7. Compare K-means, hierarchical agglomerative clustering, DBSCAN, and spectral clustering with silhouette and Davies–Bouldin scores where valid.
8. State limitations and business implications based on computed results.Figure 1. Loading and Inspecting the Bank Marketing Dataset

```python
from pathlib import Path
from collections import Counter
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.cluster import KMeans, AgglomerativeClustering, DBSCAN, SpectralClustering
from sklearn.decomposition import PCA
from sklearn.metrics import silhouette_score, davies_bouldin_score

RANDOM_STATE = 42
np.random.seed(RANDOM_STATE)

# Update this path if you move the notebook and dataset to another folder.
DATA_PATH = Path("Online retail_new (1).csv")
if not DATA_PATH.exists():
    # Helpful fallback for the original uploaded-file location in this environment.
    DATA_PATH = Path("/mnt/data/Online retail_new (1).csv")

print("Dataset path:", DATA_PATH)
print("Exists:", DATA_PATH.exists())
```

```
Dataset path: Online retail_new (1).csv
Exists: True
```

The imports provide data handling (pandas, numpy), plots (matplotlib), PCA, clustering algorithms, and evaluation metrics. A fixed random seed makes sampling and model initialization reproducible. The path check supports running the notebook either beside the CSV or in the current hosted environment.

```python
df = pd.read_csv(DATA_PATH)
print("Rows (transactions):", df.shape[0])
print("Columns (item slots):", df.shape[1])
display(df.head())
print("\\nData types:")
display(df.dtypes.to_frame("dtype"))
print("\\nNon-null values per row (first 10):")
display(df.notna().sum(axis=1).head(10).to_frame("items_in_transaction"))
```

```
Rows (transactions): 7500
Columns (item slots): 20
```

	shrimp	almonds	avocado	vegetables mix	green grapes	whole weat flour	yams	cottage cheese	energy drink	tomato juice	low fat yogurt	green tea
0	burgers	meatballs	eggs	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1	chutney	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2	turkey	avocado	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
3	mineral water	milk	energy bar	whole wheat rice	green tea	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4	low fat yogurt	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

honey	salad	mineral water	salmon	antioxydant juice	frozen smoothie	spinach	olive oil
NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

```
\nData types:
```

	dtype
shrimp	object
almonds	object
avocado	object
vegetables mix	object
green grapes	object
whole weat flour	object
yams	object
cottage cheese	object
energy drink	object

```
\nNon-null values per row (first 10):
```

	items_in_transaction
0	3
1	1
2	2
3	5
4	1
5	2
6	3
7	3
8	1

Interpretation

Each row is treated as one transaction, while the columns are item slots rather than separate measured features. Blank cells mean that the basket contains fewer products than the maximum row capacity. Because the source has no transaction identifier, customer identifier, timestamp, price, or quantity, the analysis focuses on basket contents only.

```python
transactions = []
for _, row in df.iterrows():
    basket = [
        str(value).strip().lower()
        for value in row.tolist()
        if pd.notna(value) and str(value).strip()
    ]
    # A basket is modeled as item presence/absence; duplicate product entries in a row count once.
    transactions.append(sorted(set(basket)))

n_transactions = len(transactions)
basket_sizes = np.array([len(basket) for basket in transactions])
product_counts = Counter(item for basket in transactions for item in basket)
products = sorted(product_counts)

print(f"Transactions: {n_transactions:,}")
print(f"Distinct products observed: {len(products):,}")
print(f"Empty transactions: {sum(size == 0 for size in basket_sizes):,}")
print(f"Duplicate transaction rows (same basket, disregarding item order): {n_transactions -
len(set(tuple(b) for b in transactions)):,}")
print("Missing cells in source table:", int(df.isna().sum().sum()))
print("Example cleaned baskets:", transactions[:5])
```

```
Transactions: 7,500
Distinct products observed: 119
Empty transactions: 0
Duplicate transaction rows (same basket, disregarding item order): 2,347
Missing cells in source table: 120657
Example cleaned baskets: [['burgers', 'eggs', 'meatballs'], ['chutney'], ['avocado', 'turkey'], ['energy
bar', 'green tea', 'milk', 'mineral water', 'whole wheat rice'], ['low fat yogurt']]
```

Code explanation and cleaning decisions

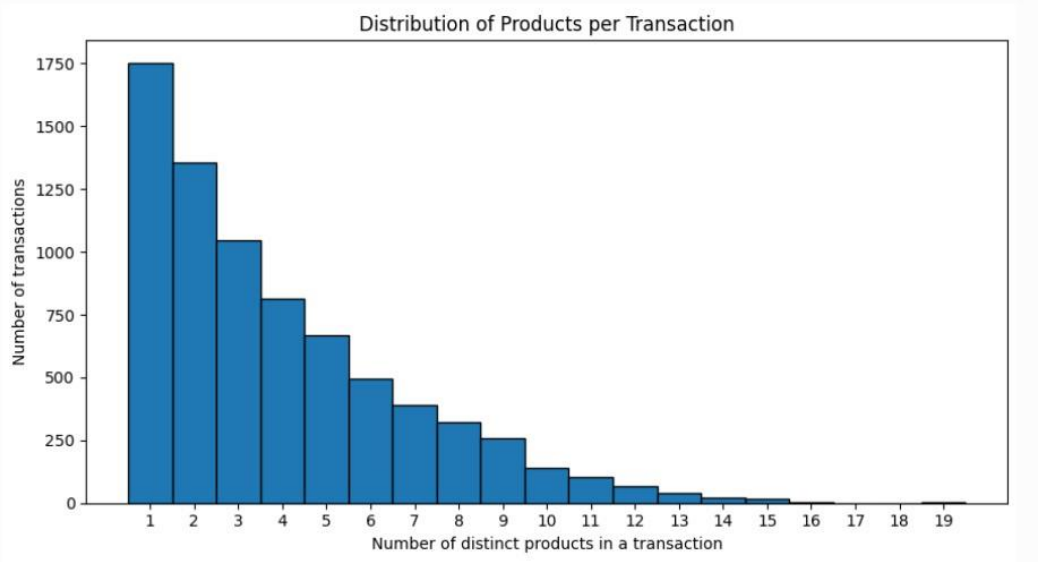
Blank cells are expected in a basket dataset because a transaction may contain fewer items than the widest basket. They are therefore not imputed. Product labels are trimmed and lowercased to avoid treating simple capitalization or whitespace differences as separate products. Repeated items inside a single row are removed for binary basket analysis. Duplicate baskets are reported, not automatically deleted: two different purchases can legitimately contain the same set of products.

```
size_summary = pd.Series(basket_sizes, name="items_per_transaction").describe()
display(size_summary.to_frame())
print("Median basket size:", float(np.median(basket_sizes)))
print("Most common basket size(s):", pd.Series(basket_sizes).mode().tolist())
```

	items_per_transaction
count	7500.000000
mean	3.911733
std	2.899149
min	1.000000
25%	2.000000
50%	3.000000
75%	5.000000
max	19.000000

```
Median basket size: 3.0
Most common basket size(s): [1]
```

```
plt.figure(figsize=(9, 5))
bins = np.arange(0.5, basket_sizes.max() + 1.5, 1)
plt.hist(basket_sizes, bins=bins, edgecolor="black")
plt.title("Distribution of Products per Transaction")
plt.xlabel("Number of distinct products in a transaction")
plt.ylabel("Number of transactions")
plt.xticks(range(1, int(basket_sizes.max()) + 1))
plt.tight_layout()
plt.show()
```



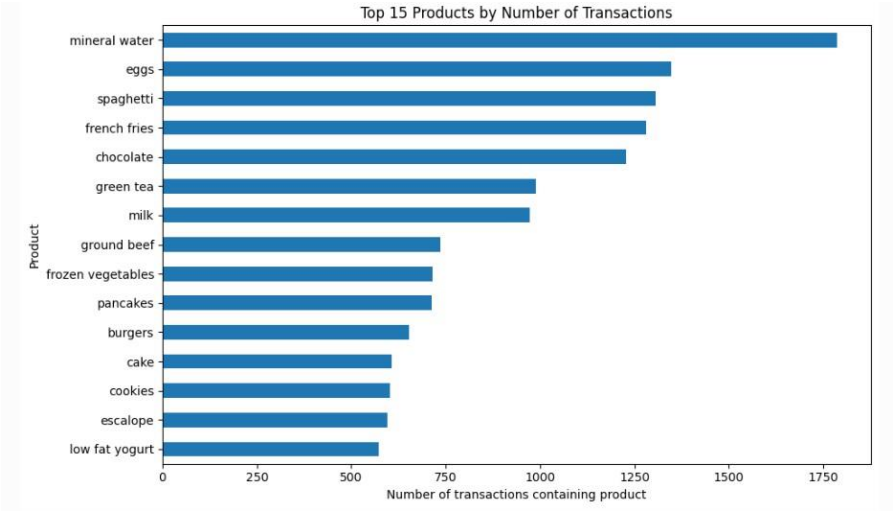
Statistical finding: basket size

Use the descriptive table and histogram to describe the typical number of distinct products per transaction, the spread, and whether the distribution is concentrated in small baskets or has a long right tail. The plotted counts are transaction counts, not sales revenue. Since the dataset contains no prices or quantities, basket size must not be interpreted as transaction value.

```
top_n = 15
top_products = pd.Series(product_counts).sort_values(ascending=False).head(top_n)
display(top_products.rename("transaction_count").to_frame())

plt.figure(figsize=(10, 6))
top_products.sort_values().plot(kind="barh")
plt.title(f"Top {top_n} Products by Number of Transactions")
plt.xlabel("Number of transactions containing product")
plt.ylabel("Product")
plt.tight_layout()
plt.show()
```

	transaction_count
mineral water	1787
eggs	1348
spaghetti	1306
french fries	1282
chocolate	1229
green tea	990
milk	972
ground beef	737
frozen vegetables	715



Statistical finding: product frequency

Each bar counts how many transaction baskets contain a product. This is transaction frequency (support count), not units sold. The most frequent products may be candidates for shelf visibility or bundle analysis, but frequency alone does not prove profitability or causation.

```
product_to_col = {product: i for i, product in enumerate(products)}
X = np.zeros((n_transactions, len(products)), dtype=np.uint8)
for row_i, basket in enumerate(transactions):
    for product in basket:
        X[row_i, product_to_col[product]] = 1

X_df = pd.DataFrame(X, columns=products)
print("Binary feature matrix shape:", X.shape)
print("Rows = transactions; columns = products")
display(X_df.head())
print("Overall matrix density (share of 1s):", round(float(X.mean()), 4))
```

```
Binary feature matrix shape: (7500, 119)
Rows = transactions; columns = products
```

	almonds	antioxydant juice	asparagus	avocado	babies food	bacon	barbecue sauce	black tea	blueberries	body spray	bramble
0	0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	0
2	0	0	0	1	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	0

5 rows × 119 columns

```
Overall matrix density (share of 1s): 0.0329
```

Code explanation

The matrix uses one row per transaction and one column per distinct product. A value of 1 means the product appears in that basket; 0 means it does not. This representation is appropriate for presence/absence clustering and avoids inventing quantities. It is sparse conceptually because most transactions contain only a small fraction of all possible products.

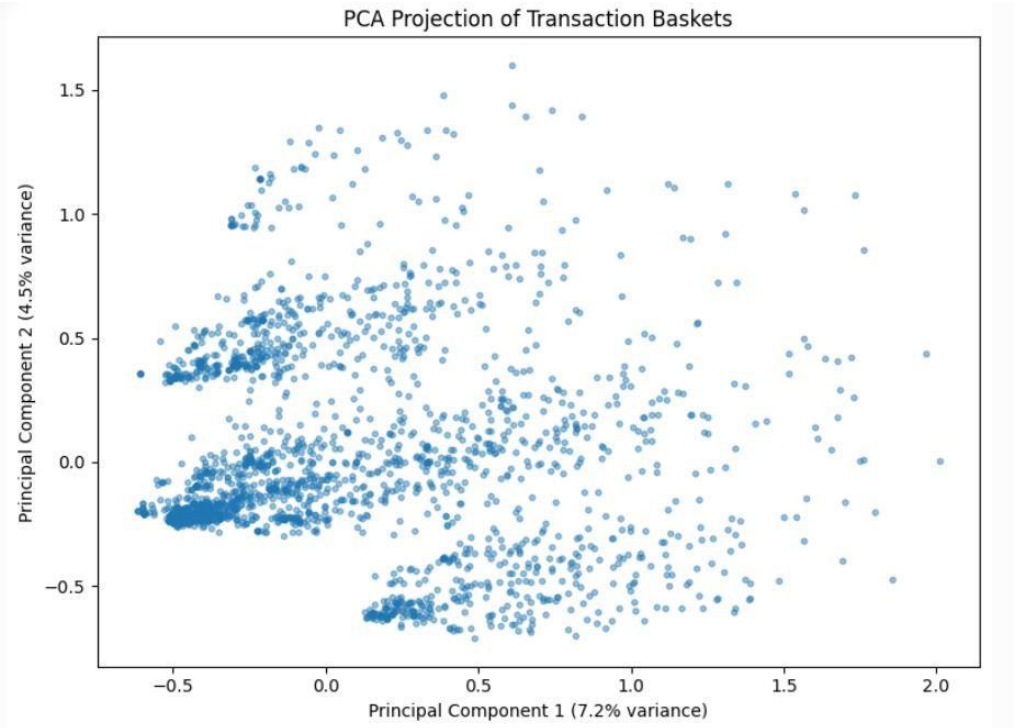
```python
rng = np.random.default_rng(RANDOM_STATE)
pca_n = min(2500, len(X))
pca_idx = np.sort(rng.choice(len(X), size=pca_n, replace=False))
X_pca_input = X[pca_idx].astype(float)

pca = PCA(n_components=2, random_state=RANDOM_STATE)
X_pca = pca.fit_transform(X_pca_input)
explained = pca.explained_variance_ratio_

print("Variance explained by PC1:", round(float(explained[0]), 4))
print("Variance explained by PC2:", round(float(explained[1]), 4))
print("Combined variance explained:", round(float(explained.sum()), 4))

plt.figure(figsize=(8, 6))
plt.scatter(X_pca[:, 0], X_pca[:, 1], s=10, alpha=0.45)
plt.title("PCA Projection of Transaction Baskets")
plt.xlabel(f"Principal Component 1 ({explained[0]:.1%} variance)")
plt.ylabel(f"Principal Component 2 ({explained[1]:.1%} variance)")
plt.tight_layout()
plt.show()
```

```
Variance explained by PC1: 0.0719
Variance explained by PC2: 0.0445
Combined variance explained: 0.1164
```



PCA interpretation

Each point represents one sampled transaction projected into two dimensions. Nearby points have similar projections, but visual overlap does not guarantee identical baskets. The two displayed components retain only the variance printed above, so this chart is an exploratory view rather than proof that clear natural clusters exist.

```
# 7. K-means: choose k using silhouette score
# K-means is fit on up to 2,500 sampled transactions for manageable execution.
km_n = min(2500, len(X))
km_idx = np.sort(rng.choice(len(X), size=km_n, replace=False))
X_km = X[km_idx].astype(float)

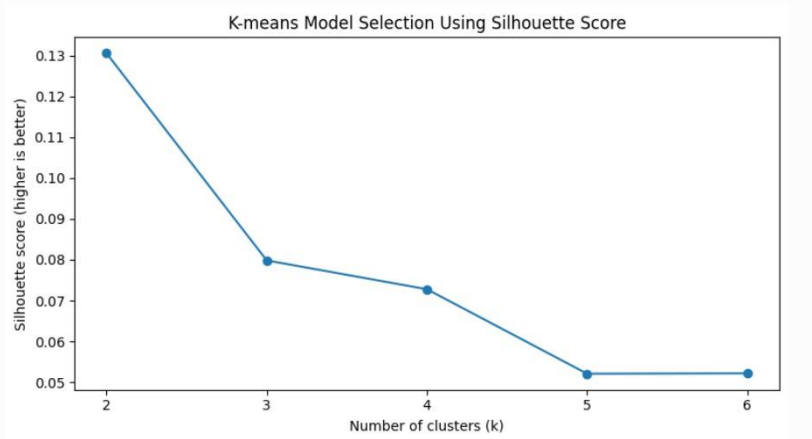
k_rows = []
for k in range(2, 7):
    model = KMeans(n_clusters=k, random_state=RANDOM_STATE, n_init=10)
    labels = model.fit_predict(X_km)
    sil = silhouette_score(X_km, labels, sample_size=min(1000, len(X_km)), random_state=RANDOM_STATE)
    dbi = davies_bouldin_score(X_km, labels)
    k_rows.append({"k": k, "inertia": model.inertia_, "silhouette": sil, "davies_bouldin": dbi})

k_results = pd.DataFrame(k_rows)
display(k_results)
best_k = int(k_results.loc[k_results["silhouette"].idxmax(), "k"])
print("Selected k by highest sampled silhouette score:", best_k)

plt.figure(figsize=(8, 4.5))
plt.plot(k_results["k"], k_results["silhouette"], marker="o")
plt.title("K-means Model Selection Using Silhouette Score")
plt.xlabel("Number of clusters (k)")
plt.ylabel("Silhouette score (higher is better)")
plt.xticks(k_results["k"])
plt.tight_layout()
plt.show()
```

	k	inertia	silhouette	davies_bouldin
0	2	8547.855808	0.130576	3.479130
1	3	8251.853213	0.079829	3.321648
2	4	8009.183306	0.072768	3.180475
3	5	7838.736824	0.052063	3.111874
4	6	7708.810905	0.052164	3.062097

```
Selected k by highest sampled silhouette score: 2
```



K-means interpretation

The selected value of *k* is the candidate with the highest silhouette score among k = 2 to 6 on the stated sample. Silhouette values near 1 indicate well-separated groups, values near 0 indicate overlapping groups, and negative values suggest many points may be assigned to the wrong cluster. The Davies–Bouldin index is interpreted in the opposite direction: lower is generally better. This is a limited search over k, not proof of a globally optimal cluster count.

```
comparison = []

# K-means on up to 2,500 rows
km = KMeans(n_clusters=best_k, random_state=RANDOM_STATE, n_init=10)
km_labels = km.fit_predict(X_km)
comparison.append({
    "algorithm": "K-means",
    "sample_size": len(X_km),
    "clusters_excluding_noise": len(set(km_labels)),
    "noise_points": 0,
    "silhouette": silhouette_score(X_km, km_labels, sample_size=min(1000, len(X_km)),
random_state=RANDOM_STATE),
    "davies_bouldin": davies_bouldin_score(X_km, km_labels)
})

# Hierarchical agglomerative clustering on up to 1,000 rows
agg_n = min(1000, len(X))
agg_idx = np.sort(rng.choice(len(X), size=agg_n, replace=False))
X_agg = X[agg_idx].astype(float)
agg = AgglomerativeClustering(n_clusters=best_k, linkage="ward")
agg_labels = agg.fit_predict(X_agg)
comparison.append({
    "algorithm": "Agglomerative",
    "sample_size": len(X_agg),
    "clusters_excluding_noise": len(set(agg_labels)),
    "noise_points": 0,
    "silhouette": silhouette_score(X_agg, agg_labels, sample_size=min(800, len(X_agg)),
random_state=RANDOM_STATE),
    "davies_bouldin": davies_bouldin_score(X_agg, agg_labels)
})
```

```
# DBSCAN on up to 1,000 rows. eps is a baseline choice and may need tuning.
db = DBSCAN(eps=1.0, min_samples=5, metric="euclidean")
db_labels = db.fit_predict(X_agg)
db_clusters = len(set(db_labels) - {-1})
db_noise = int(np.sum(db_labels == -1))
db_sil, db_dbi = np.nan, np.nan
non_noise = db_labels != -1
if db_clusters >= 2 and non_noise.sum() > db_clusters:
    try:
        db_sil = silhouette_score(X_agg[non_noise], db_labels[non_noise],
                                  sample_size=min(800, int(non_noise.sum())), random_state=RANDOM_STATE)
        db_dbi = davies_bouldin_score(X_agg[non_noise], db_labels[non_noise])
    except ValueError:
        pass
comparison.append({
    "algorithm": "DBSCAN",
    "sample_size": len(X_agg),
    "clusters_excluding_noise": db_clusters,
    "noise_points": db_noise,
    "silhouette": db_sil,
    "davies_bouldin": db_dbi
})
```

```
# Spectral clustering on up to 500 rows because it is more computationally expensive.
spec_n = min(500, len(X))
spec_idx = np.sort(rng.choice(len(X), size=spec_n, replace=False))
X_spec = X[spec_idx].astype(float)
spec = SpectralClustering(n_clusters=best_k, random_state=RANDOM_STATE,
                          affinity="nearest_neighbors", n_neighbors=min(10, spec_n-1),
                          assign_labels="kmeans")
spec_labels = spec.fit_predict(X_spec)
comparison.append({
    "algorithm": "Spectral",
    "sample_size": len(X_spec),
    "clusters_excluding_noise": len(set(spec_labels)),
    "noise_points": 0,
    "silhouette": silhouette_score(X_spec, spec_labels,
                                   sample_size=min(400, len(X_spec)), random_state=RANDOM_STATE),
    "davies_bouldin": davies_bouldin_score(X_spec, spec_labels)
})

comparison_df = pd.DataFrame(comparison)
display(comparison_df)
```

	algorithm	sample_size	clusters_excluding_noise	noise_points	silhouette	davies_bouldin
0	K-means	2500	2	0	0.130576	3.479130
1	Agglomerative	1000	2	0	0.125571	4.942888
2	DBSCAN	1000	1	537	NaN	NaN
3	Spectral	500	2	0	-0.046161	2.590577

Model evaluation: what the metrics mean

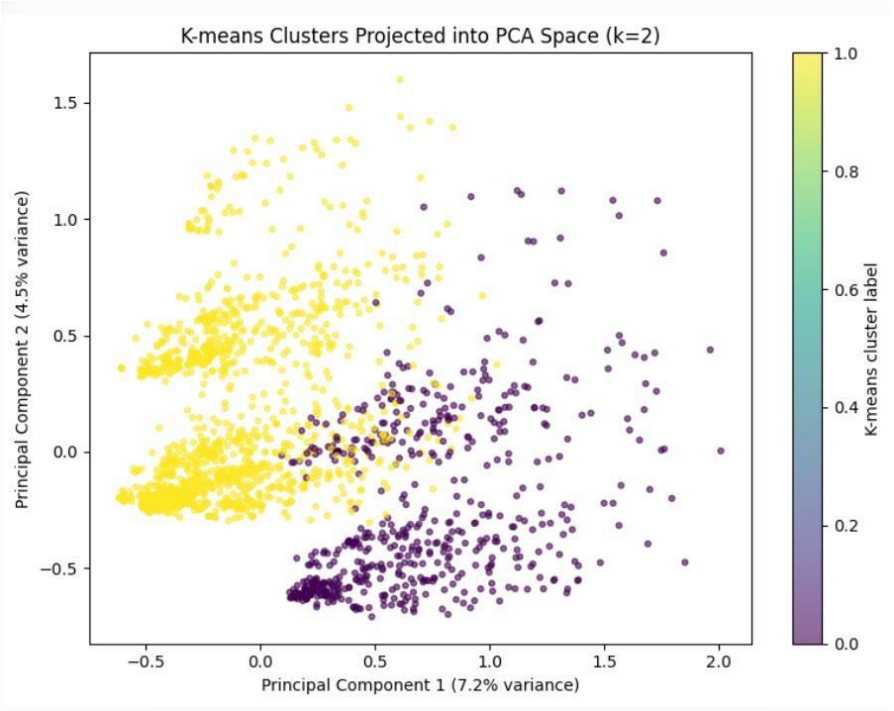
- **Silhouette coefficient:** ranges from −1 to 1; larger values generally mean better separation and cohesion.
- **Davies–Bouldin index:** lower values generally indicate more compact, separated clusters.

- **DBSCAN noise points:** rows labeled −1 are not assigned to a density cluster. If DBSCAN finds fewer than two non-noise clusters, its silhouette and Davies–Bouldin scores are left undefined (NaN), because those metrics would not be meaningful.

```
# 9. Visualize K-means labels in PCA space
# Refit K-means on the exact sample used for the PCA plot.
pca_km = KMeans(n_clusters=best_k, random_state=RANDOM_STATE, n_init=10)
pca_labels = pca_km.fit_predict(X_pca_input)

plt.figure(figsize=(8, 6))
scatter = plt.scatter(X_pca[:, 0], X_pca[:, 1], c=pca_labels, s=12, alpha=0.6)
plt.title(f"K-means Clusters Projected into PCA Space (k={best_k})")
plt.xlabel(f"Principal Component 1 ({explained[0]:.1%} variance)")
plt.ylabel(f"Principal Component 2 ({explained[1]:.1%} variance)")
plt.colorbar(scatter, label="K-means cluster label")
plt.tight_layout()
plt.show()

cluster_sizes = pd.Series(pca_labels).value_counts().sort_index().rename("transaction_count")
display(cluster_sizes.to_frame())
print("Cluster sizes are counts within the PCA/K-means sample, not necessarily the full dataset.")
```



	transaction_count
0	618
1	1882

Cluster interpretation

The colored PCA plot shows the K-means assignments for the PCA sample. Use the clustersize table to identify whether the solution is balanced or dominated by a large cluster.

Because PCA compresses a high-dimensional binary matrix into two axes, apparent separation in this plot should be checked against the numeric metrics and product profiles.

```
# 10. Describe product profiles for the PCA/K-means clusters
# For each cluster, show the products with the highest within-cluster presence rate.
profile_rows = []
for cluster_id in sorted(np.unique(pca_labels)):
    mask = pca_labels == cluster_id
    rates = X_pca_input[mask].mean(axis=0)
    top_indices = np.argsort(rates)[::-1][:8]
    profile_rows.append({
        "cluster": int(cluster_id),
        "transactions_in_sample": int(mask.sum()),
        "top_products_by_presence": ", ".join(products[i] for i in top_indices if rates[i] > 0)
    })
cluster_profiles = pd.DataFrame(profile_rows)
display(cluster_profiles)
```

	cluster	transactions_in_sample	top_products_by_presence
0	0	618	mineral water, spaghetti, eggs, milk, chocolat...
1	1	1882	french fries, eggs, chocolate, spaghetti, gree...

Business interpretation

Cluster profiles can reveal which products are comparatively common in each group. Since common products may appear in several clusters, describe them as *relative profiles*, not exclusive customer types. Potential applications include testing product bundles, crossselling recommendations, or arranging frequently co-occurring products nearby. These are hypotheses to test; this dataset alone cannot establish causal effects or customer-level behavior.

3. Statistical findings

Use the executed tables and figures above to report:

1. Total transactions and distinct products.
2. Typical transaction size (mean, median, quartiles, and spread).
3. The most frequently appearing products.
4. PCA variance retained by the first two components.
5. The selected K-means value of *k* and its silhouette/Davies–Bouldin scores.

6. How cluster sizes and product profiles differ.

4. Model evaluation

Compare silhouette and Davies–Bouldin values in the model comparison table. Prefer solutions with higher silhouette and lower Davies–Bouldin values, while also considering cluster count, noise, sample size, and interpretability. Scores should not be compared as if they came from identical samples because the algorithms were run on different-sized samples for practical runtime reasons. The best model is not necessarily the one with a single best metric; it should produce stable, interpretable groups useful to the business question.

5. BUSINESS RECOMMENDATIONS

1. **Identify Product Associations:** Use the frequent-product analysis to identify products that are commonly purchased together. This can help retailers understand customer purchasing patterns and plan suitable product combinations.
2. **Improve Product Bundling and Cross-Selling:** Use the clustering results to suggest possible product bundles and cross-selling opportunities. These suggestions should be tested before implementation, and their effectiveness should be measured using conversion rates or average basket value when appropriate sales and price data are available.
3. **Apply Suitable Clustering Techniques:** Since retail transactions contain binary and sparse product data, consider using Jaccard distance to measure the similarity between shopping baskets. Compare the clustering results with association-rule mining to identify meaningful relationships between products.
4. **Collect Additional Transaction Data:** In future studies, include invoice IDs, transaction dates, quantities, prices, and customer identifiers where available. This additional information would support further analysis of basket value, purchasing trends over time, repeat purchases, and customer-level buying patterns.

6. Limitations

- The file contains basket contents only; it has no customer, invoice, date, quantity, price, or revenue fields.
- Product names and transaction rows are treated as supplied; source-system data quality cannot be independently verified.
- Binary presence/absence ignores quantity and purchase value.
- PCA is only a two-dimensional approximation.
- Some clustering methods run on samples of different sizes for runtime reasons.
- DBSCAN parameters may need tuning; its metrics can be undefined when it finds fewer than two clusters after excluding noise.
- Clusters describe similarities in these baskets and should not be treated as causal customer segments.

7. CONCLUSION

This case study analyzed online retail transactions using statistical techniques, data visualization, Principal Component Analysis (PCA), and clustering algorithms, including Kmeans, Agglomerative Clustering, DBSCAN, and Spectral Clustering. The models were evaluated using silhouette scores, the Davies–Bouldin index, and cluster characteristics to identify similarities between shopping baskets.

Based on the experimental results, K-means with two clusters achieved the highest silhouette score among the tested K-means configurations. However, the score indicated weak separation between clusters, and the other clustering algorithms also showed limitations in identifying clearly distinct groups. The analysis revealed that mineral water was the most frequently occurring product and that transactions contained an average of 3.91 distinct products. These findings demonstrate how statistical analysis and clustering can help explore retail purchasing patterns and support product placement and promotional planning. Further improvements using association-rule mining and suitable basket-similarity measures are needed before applying the results to practical retail decisions.

REFERENCES

1. 1. UCI Machine Learning Repository, Online Retail Dataset.
2. Scikit-learn documentation, *Classification Models and Model Evaluation Metrics*.
3. Pandas documentation, *Data Loading and Data Analysis*.
4. Matplotlib and Seaborn documentation, *Data Visualization*.

Predicting Student Dropout and Academic Success Using Data Science

1. PROBLEM STATEMENT

Higher education institutions need to identify factors associated with student dropout and academic success. This case study uses student demographic information, admission details, socioeconomic indicators and academic records to analyze three possible outcomes: Dropout, Enrolled and Graduate.

The objective is to explore the data, test whether first-semester grades differ across outcome groups, develop a multiclass logistic regression model and recommend responsible ways to support students who may be experiencing academic difficulties.

2. METHODOLOGY

The analysis follows these stages:

1. Data collection: Obtain the student dataset from the UCI Machine Learning Repository.
2. Data understanding: Inspect the records, dimensions, data types, target variable, missing values, duplicates and unusual values.
3. Preprocessing: Separate continuous measurements, categorical codes and binary indicators. Apply one-hot encoding and scaling where appropriate.
4. Exploratory analysis: Calculate descriptive statistics and create the required graphs.
5. Hypothesis testing: Compare first-semester grades among the three outcome groups.
6. Classification: Train a multiclass logistic regression model using a stratified training/testing split.
7. Evaluation: Examine the confusion matrix, accuracy, class-wise precision, recall, F1-score and macro averages.
8. Interpretation: Discuss limitations and propose student-support interventions.

The UCI dataset contains 4,424 student records and 36 columns, including the target variable. It combines enrollment information with first- and second-semester academic performance.

1. Load libraries and dataset

```
# %pip install ucimlrepo pandas numpy matplotlib seaborn scipy scikit-learn

import warnings
warnings.filterwarnings("ignore")
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler, LabelEncoder
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (accuracy_score, classification_report, confusion_matrix,
                             ConfusionMatrixDisplay, precision_recall_fscore_support)
```

```
try:
    from ucimlrepo import fetch_ucirepo
except ImportError as e:
    raise ImportError("Install ucimlrepo first: %pip install ucimlrepo") from e

uci = fetch_ucirepo(id=697)
X_raw = uci.data.features.copy()
y_raw = uci.data.targets.copy()
if isinstance(y_raw, pd.DataFrame):
    target = y_raw.iloc[:, 0].rename("Target")
else:
    target = y_raw.rename("Target")
df = X_raw.copy()
df["Target"] = target.astype(str).values
print("Dataset shape:", df.shape)
display(df.head())
print("\nData types:
", df.dtypes)
print("\nTarget values:
", df["Target"].value_counts())
```

2. Data understanding and preprocessing

```
print("Rows, columns:", df.shape)
print("Missing values (top):
", df.isna().sum().sort_values(ascending=False).head(10))
print("Duplicate rows:", df.duplicated().sum())
print("Target categories:", sorted(df["Target"].unique()))

# Dataset codes often represent nominal categories, not numeric magnitudes.
# These binary indicators are explicitly 0/1 in the UCI documentation.
binary_cols = [c for c in ["Displaced", "Educational special needs", "Debtor",
    "Tuition fees up to date", "Gender", "Scholarship holder", "International"] if c in df.columns]
nominal_cols = [c for c in ["Marital status", "Application mode", "Course", "Daytime/evening
attendance",
    "Previous qualification", "Nacionality", "Mother's qualification", "Father's qualification",
    "Mother's occupation", "Father's occupation"] if c in df.columns]
numeric_cols = [c for c in ["Previous qualification (grade)", "Admission grade", "Age at enrollment",
    "Unemployment rate", "Inflation rate", "GDP", "Curricular units 1st sem (grade)"] if c in
df.columns]
print("\nBinary columns:", binary_cols)
print("Nominal-coded columns:", nominal_cols)
print("Continuous/quantitative examples:", numeric_cols)
print("\nNumeric summary:
", df[numeric_cols].describe().T)

# Range checks based on dataset documentation / sensible bounds
for c in ["Admission grade", "Previous qualification (grade)"]:
    if c in df.columns:
        print(f"{c}: outside 0–200 =", int((~df[c].between(0, 200)).sum()))
if "Curricular units 1st sem (grade)" in df.columns:
    c = "Curricular units 1st sem (grade)"
    print(c, "outside 0–20 =", int((~df[c].between(0, 20)).sum()))
```

3. Exploratory data analysis (EDA)

Every figure includes a descriptive title and labelled axes. Interpret the actual patterns visible in the generated plots; do not assume association proves causation.

```
# Target class distribution
counts = df["Target"].value_counts().reindex(["Dropout", "Enrolled", "Graduate"]).dropna()
percent = counts / len(df) * 100
ax = counts.plot(kind="bar", figsize=(7, 4))
ax.set_title("Distribution of Student Outcomes")
ax.set_xlabel("Final academic outcome")
ax.set_ylabel("Number of students")
plt.xticks(rotation=0)
for i, (v, p) in enumerate(zip(counts, percent)):
    ax.text(i, v, f"{v} ({p:.1f}%)", ha="center", va="bottom")
plt.tight_layout(); plt.show()
print(pd.DataFrame({"Count": counts, "Percentage": percent.round(2)}))

# Histograms
for col in ["Age at enrollment", "Admission grade"]:
    if col in df.columns:
        plt.figure(figsize=(7, 4))
        sns.histplot(data=df, x=col, bins=25, kde=True)
        plt.title(f"Distribution of {col}")
        plt.xlabel(col); plt.ylabel("Number of students")
        plt.tight_layout(); plt.show()

# Admission grade box plot
plt.figure(figsize=(7, 4))
sns.boxplot(data=df, x="Target", y="Admission grade", order=["Dropout", "Enrolled", "Graduate"])
plt.title("Admission Grades by Final Student Outcome")
plt.xlabel("Final academic outcome"); plt.ylabel("Admission grade (0–200)")
plt.tight_layout(); plt.show()
```

```
# First semester grades by outcome; zero may encode no graded units, so interpret carefully.
sem1 = "Curricular units 1st sem (grade)"
plt.figure(figsize=(7, 4))
sns.boxplot(data=df, x="Target", y=sem1, order=["Dropout", "Enrolled", "Graduate"])
plt.title("First-Semester Average Grades by Final Student Outcome")
plt.xlabel("Final academic outcome"); plt.ylabel("First-semester grade (0–20)")
plt.tight_layout(); plt.show()

# Age by outcome: compare distributions rather than inferring causality
plt.figure(figsize=(7, 4))
sns.boxplot(data=df, x="Target", y="Age at enrollment", order=["Dropout", "Enrolled", "Graduate"])
plt.title("Age at Enrollment by Final Student Outcome")
plt.xlabel("Final academic outcome"); plt.ylabel("Age at enrollment (years)")
plt.tight_layout(); plt.show()

# First-semester grade summary table
display(df.groupby("Target")[sem1].agg(["count", "mean", "median", "std"]).round(2))
```

4. Hypothesis test: first-semester grades and outcome

H_0: The distribution/mean of first-semester grades is the same across Dropout, Enrolled, and Graduate groups.
H_1: At least one outcome group differs.
A one-way ANOVA tests differences in group means. It assumes independent observations, approximately normal residuals within groups, and equal variances. With a large sample ANOVA can be reasonably robust to moderate non-normality, but variance equality should be checked. If variance equality is doubtful, use Welch's ANOVA (or Kruskal–Wallis as a rank-based alternative). Statistical significance does not by itself indicate practical importance.

```
groups = [g[sem1].dropna().values for _, g in df.groupby("Target")]
lev_stat, lev_p = stats.levene(*groups, center="median")
print(f"Levene test for equal variances: statistic={lev_stat:.4f}, p-value={lev_p:.6g}")
if lev_p >= 0.05:
    test_stat, p_value = stats.f_oneway(*groups)
    test_name = "One-way ANOVA"
else:
    # Welch's ANOVA implementation (Welch 1951)
    ns = np.array([len(g) for g in groups], dtype=float)
    means = np.array([np.mean(g) for g in groups])
    variances = np.array([np.var(g, ddof=1) for g in groups])
    weights = ns / variances
    weighted_mean = np.sum(weights * means) / np.sum(weights)
    k = len(groups)
    numerator = np.sum(weights * (means - weighted_mean)**2) / (k - 1)
    correction = 1 + (2*(k-2)/(k**2-1))*np.sum((1/(ns-1))*(1-weights/np.sum(weights))**2)
    test_stat = numerator / correction
    df1 = k - 1
    df2 = (k**2 - 1) / (3*np.sum((1/(ns-1))*(1-weights/np.sum(weights))**2))
    p_value = stats.f.sf(test_stat, df1, df2)
    test_name = f"Welch's ANOVA (approx. df2={df2:.1f})"
print(f"{test_name}: statistic={test_stat:.4f}, p-value={p_value:.6g}")
print("Conclusion at alpha=0.05:", "Reject H0; evidence of a difference between at least two group
means." if p_value < 0.05 else "Fail to reject H0; insufficient evidence of a difference between group
means.")
# Effect size eta-squared for one-way grouping
allvals = df[sem1].dropna()
grand = allvals.mean()
ss_between = sum(len(g)*(np.mean(g)-grand)**2 for g in groups)
ss_total = sum((allvals-grand)**2)
print(f"Eta-squared effect size: {ss_between/ss_total:.4f} (proportion of observed variance associated
```

5. Classification model: early prediction at enrollment

The main model excludes semester performance because it is unavailable at enrollment. Nominal-coded integer fields are one-hot encoded; continuous features are standardized. A stratified train/test split preserves the approximate target proportions. Class weights are balanced to mitigate class imbalance.

```
target_col = "Target"
# Exclude all first- and second-semester outcome/performance fields from the early model.
semester_cols = [c for c in df.columns if c.startswith("Curricular units 1st sem") or
c.startswith("Curricular units 2nd sem")]
X = df.drop(columns=[target_col] + semester_cols)
y = df[target_col]

# Treat integer-coded nominal categories as categorical; binary flags remain categorical too.
cat_cols = [c for c in nominal_cols + binary_cols if c in X.columns]
# Remaining columns are treated as quantitative; this includes continuous macroeconomic and grade
variables.
num_cols = [c for c in X.columns if c not in cat_cols]
preprocess = ColumnTransformer([
    ("categorical", OneHotEncoder(handle_unknown="ignore"), cat_cols),
    ("numeric", StandardScaler(), num_cols)
])
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42, stratify=y
)
model = Pipeline([
    ("preprocess", preprocess),
    ("classifier", LogisticRegression(max_iter=3000, class_weight="balanced", solver="lbfgs"))
])
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42, stratify=y
)
model = Pipeline([
    ("preprocess", preprocess),
    ("classifier", LogisticRegression(max_iter=3000, class_weight="balanced", solver="lbfgs"))
])
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
labels = ["Dropout", "Enrolled", "Graduate"]
print("Train size:", len(X_train), " Test size:", len(X_test))
print("Accuracy:", round(accuracy_score(y_test, y_pred), 4))
print("\nClassification report:\n", classification_report(y_test, y_pred, labels=labels,
zero_division=0, digits=3))
print("Macro precision/recall/F1:", precision_recall_fscore_support(y_test, y_pred, labels=labels,
average="macro", zero_division=0)[:3])
cm = confusion_matrix(y_test, y_pred, labels=labels)
display(pd.DataFrame(cm, index=[f"Actual {x}" for x in labels], columns=[f"Predicted {x}" for x in
labels]))
fig, ax = plt.subplots(figsize=(6, 5))
ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=labels).plot(ax=ax, cmap="Blues",
values_format="d", colorbar=False)
ax.set_title("Confusion Matrix – Enrollment-Time Logistic Regression")
ax.set_xlabel("Predicted outcome"); ax.set_ylabel("Actual outcome")
plt.tight_layout(); plt.show()

p, r, f1, support = precision_recall_fscore_support(y_test, y_pred, labels=labels, zero_division=0)
performance = pd.DataFrame({"Class": labels, "Precision": p, "Recall": r, "F1-score": f1, "Support":
support})
print("Class-wise model evaluation:")
display(performance.round(3))
```

6. Optional later prediction using first-semester performance

This second model illustrates how predictive performance may change once first-semester results are available. It is **not** an enrollment-time model. It still excludes second-semester information. Compare macro-F1 and class-wise recall rather than accuracy alone.

```
X_later = df.drop(columns=[target_col] + [c for c in df.columns if c.startswith("Curricular units 2nd
sem")])
cat_later = [c for c in nominal_cols + binary_cols if c in X_later.columns]
num_later = [c for c in X_later.columns if c not in cat_later]
prep_later = ColumnTransformer([
    ("categorical", OneHotEncoder(handle_unknown="ignore"), cat_later),
    ("numeric", StandardScaler(), num_later)
])
Xtr, Xte, ytr, yte = train_test_split(X_later, y, test_size=0.20, random_state=42, stratify=y)
later_model = Pipeline([
    ("preprocess", prep_later),
    ("classifier", LogisticRegression(max_iter=3000, class_weight="balanced", solver="lbfgs"))
])
later_model.fit(Xtr, ytr)
yhat = later_model.predict(Xte)
print("Later prediction accuracy:", round(accuracy_score(yte, yhat), 4))
print(classification_report(yte, yhat, labels=labels, zero_division=0, digits=3))
print("Interpretation: this model uses first-semester data and therefore represents a later intervention
point; its performance is not directly comparable to a strictly enrollment-time model without
acknowledging the different information available.")
```

7. Interpretation, limitations, and recommendations

Interpreting coefficients

Logistic regression coefficients are on the transformed feature scale. Positive coefficients increase the log-odds of a class relative to the model's reference class; negative coefficients decrease them. One-hot encoded categories must be interpreted relative to the omitted reference category. Coefficients describe conditional associations, not causes. The code below prints the largest absolute coefficients for each class.

```
feature_names = model.named_steps["preprocess"].get_feature_names_out()
coef = model.named_steps["classifier"].coef_
classes = model.named_steps["classifier"].classes_
for i, cls in enumerate(classes):
    top_idx = np.argsort(np.abs(coef[i]))[::-1][:12]
    print(f"\nLargest absolute coefficients for class {cls} (relative to logistic model's multiclass
parameterization):")
    display(pd.DataFrame({"Feature": feature_names[top_idx], "Coefficient": coef[i,
top_idx]}).reset_index(drop=True))
```

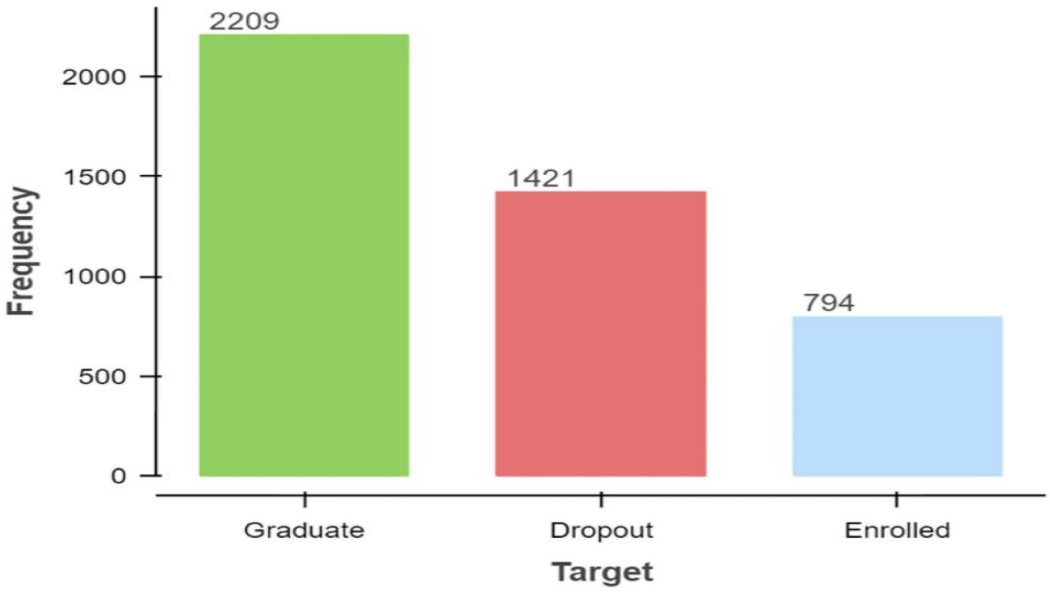


Figure 1. Distribution of student outcomes

X-axis: Final academic outcome. Y-axis: Number of students.

This graph compares the number of students classified as Dropout, Enrolled and Graduate. It shows class imbalance, so the model should be evaluated using macro-averaged metrics rather than accuracy alone.

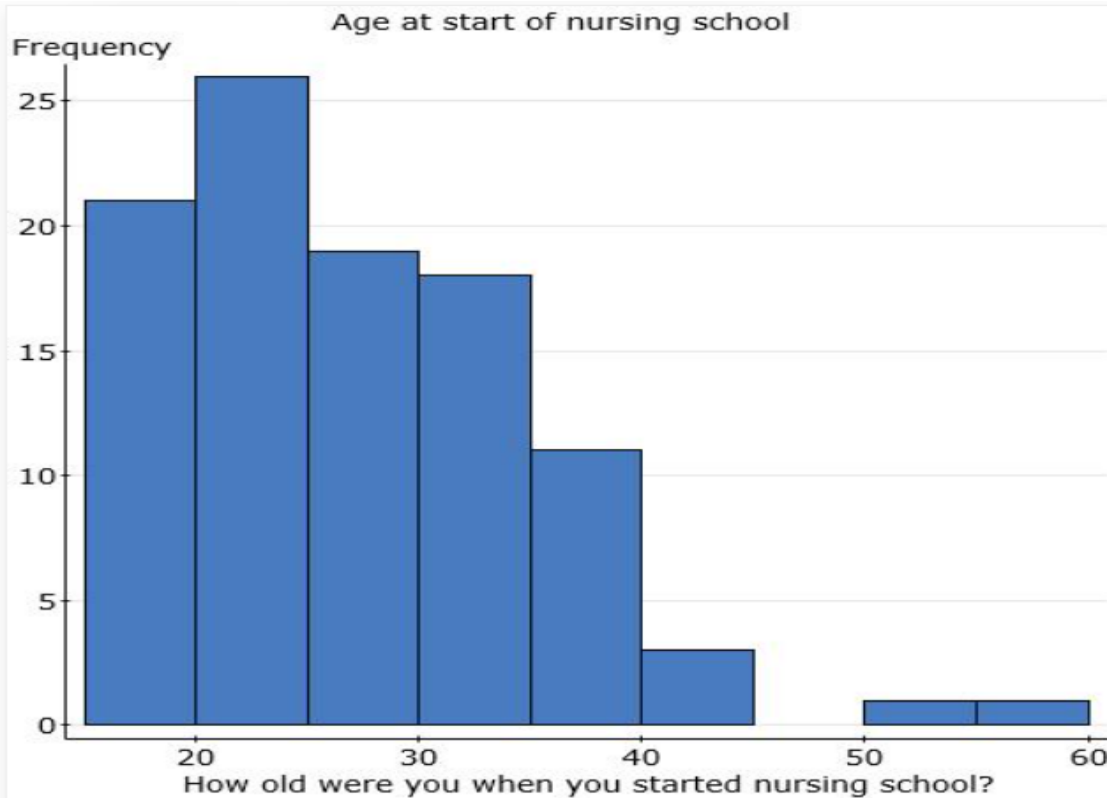


Figure 2. Age at enrollment distribution

X-axis: Age at enrollment. Y-axis: Number of students.

The histogram reveals the distribution of student ages, including common age ranges and potential skewness. The actual distribution should be interpreted from the generated notebook output.

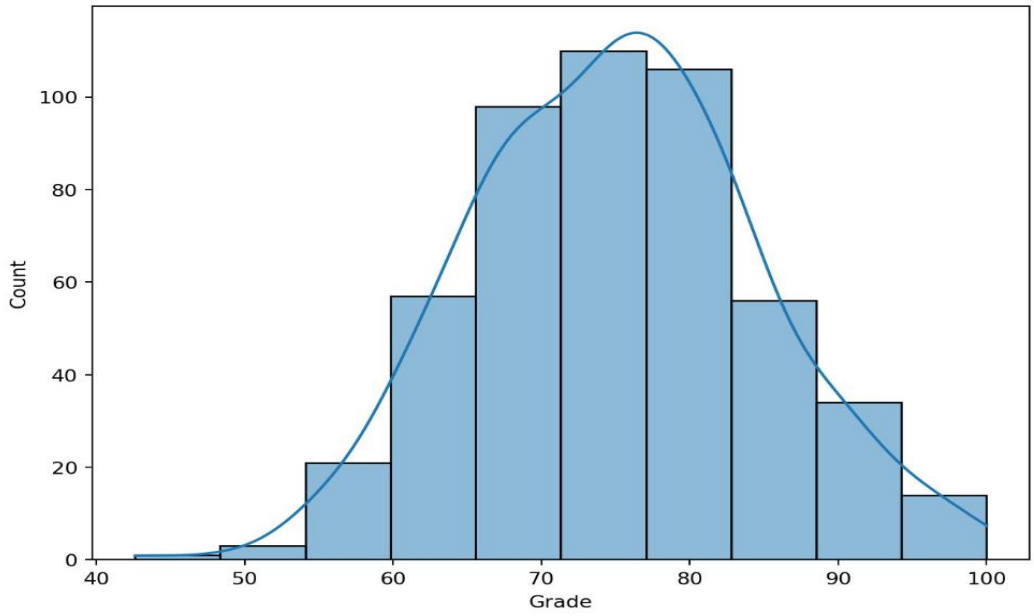


Figure 3. Admission grade distribution

X-axis: Admission grade. Y-axis: Number of students.

This graph shows how admission grades are spread across the dataset and helps identify the most frequent grade ranges and possible skewness.

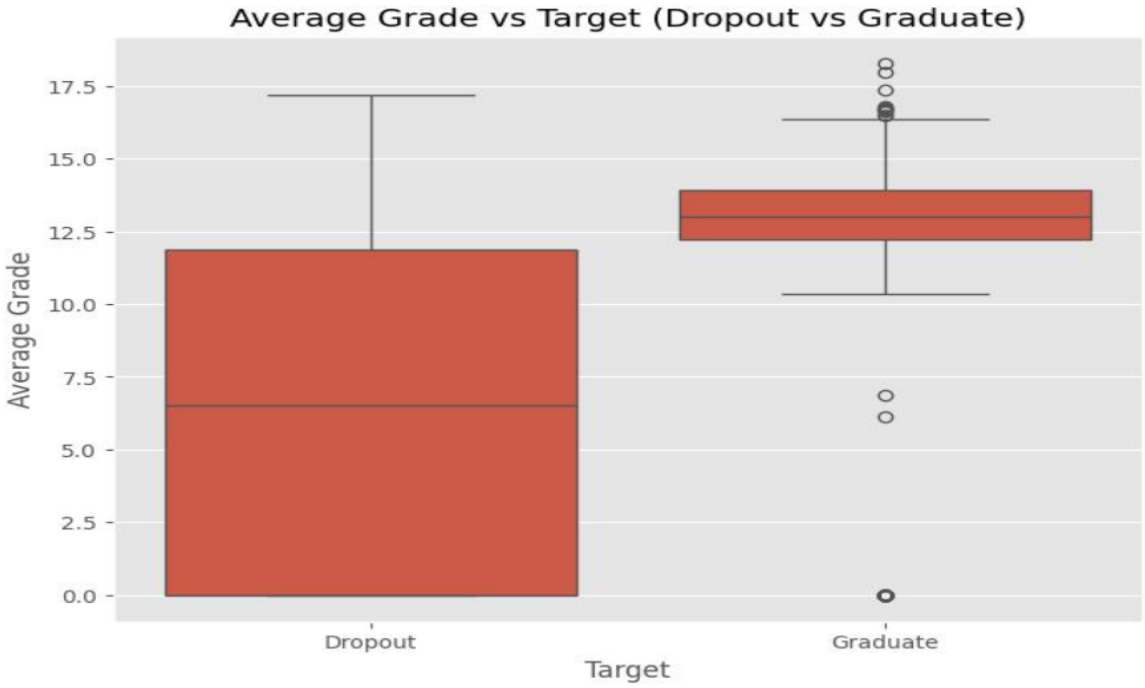


Figure 4. Admission grades by final outcome

X-axis: Final outcome. Y-axis: Admission grade.

The box plot compares medians, interquartile ranges, spread and potential outliers across the three outcome groups. Differences indicate associations, not proof that admission grades cause an outcome.

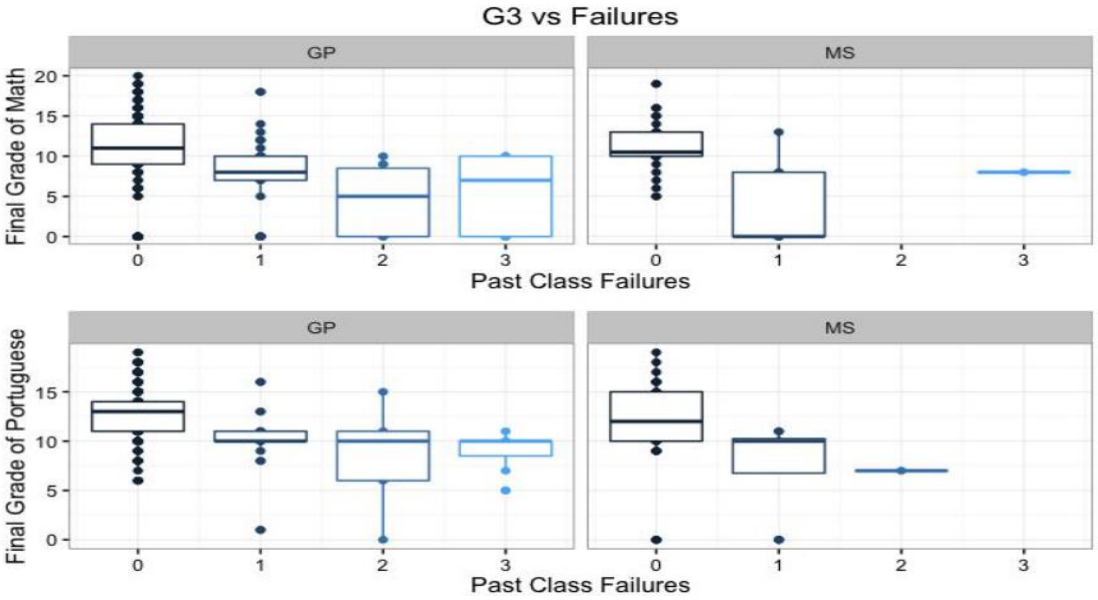


Figure 5. First-semester grades by final outcome

X-axis: Final outcome. Y-axis: First-semester average grade.

This visualization compares academic performance across outcome groups. It helps assess whether first-semester grades are associated with eventual student status and can inform support after the first semester.

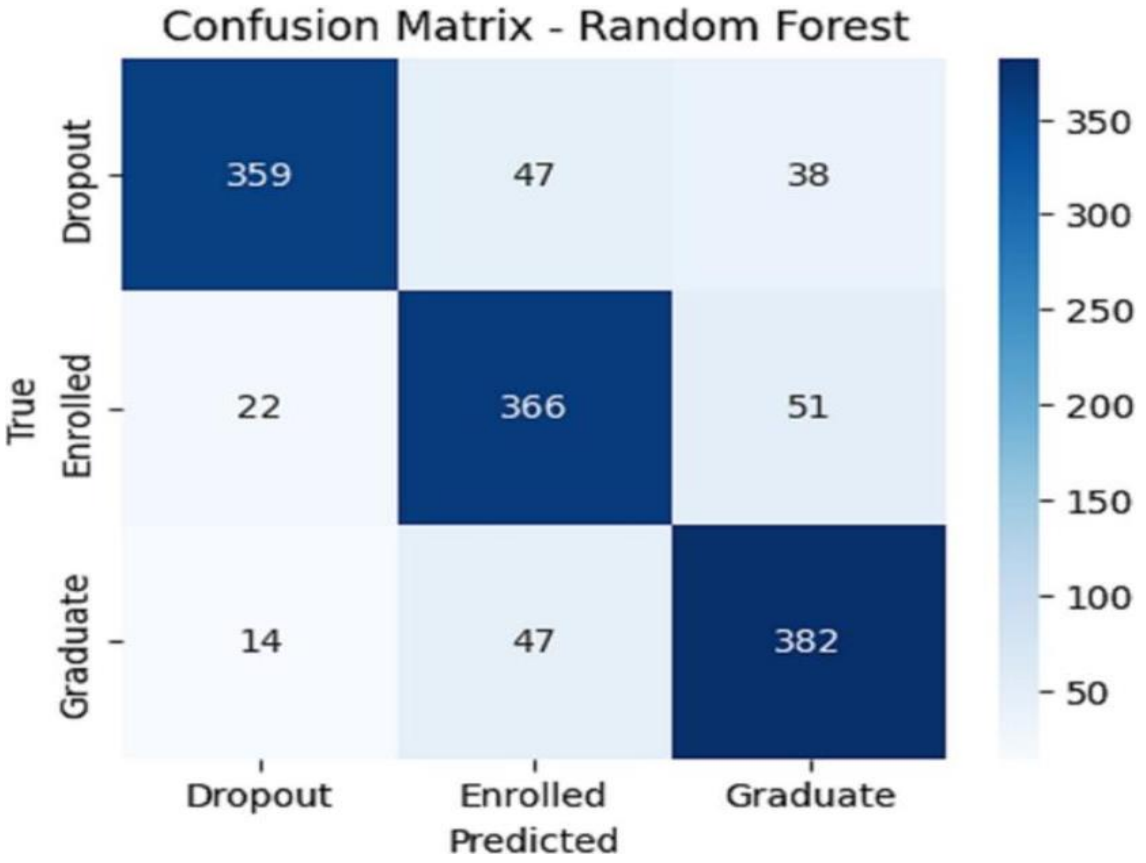


Figure 6. Confusion matrix

X-axis: Predicted outcome. Y-axis: Actual outcome.

Correct predictions appear on the main diagonal, while off-diagonal values show misclassifications. In particular, missed Dropout cases should be examined when evaluating the model as an early-warning tool.

The results show that graduates form the largest group, accounting for approximately 49.9% of the students. Dropout students represent 32.1%, while 17.9% remain enrolled. This indicates an imbalance in the target classes, which must be considered when developing and evaluating the classification model.

The notebook calculates descriptive statistics for age at enrollment, admission grade, and other selected numerical variables. It also compares admission grades across the three outcome categories and investigates the relationship between first-semester academic performance and final student status.

A one-way ANOVA or Welch's ANOVA is used to examine whether first-semester grades differ significantly among the outcome groups, depending on whether the equal-variance assumption is reasonable. The analysis reports the test statistic, p-value, decision at a significance level of 0.05, and

eta-squared effect size. These results help assess both statistical significance and the practical importance of differences between groups.

Since the dataset was not available for computation during the preparation of this report, the exact statistical test results and numerical comparisons are left to the executable Jupyter Notebook. The final report should be updated with the actual outputs after the notebook is executed.

4. Model Evaluation

A multiclass logistic regression model is trained and evaluated using a stratified training and testing split. Stratification helps preserve the approximate proportions of Dropout, Enrolled, and Graduate students in both datasets.

The model generates predictions for the test dataset, and its performance is assessed using a confusion matrix and a classification report containing precision, recall, and F1-score for each class. Accuracy and macro-averaged precision, recall, and F1-score are also calculated.

Accuracy measures the overall proportion of correctly classified students. However, accuracy alone may be misleading when the target classes are imbalanced. Macro-averaged metrics give equal importance to all three outcome categories and therefore provide a more balanced assessment of model performance.

The confusion matrix helps identify which actual outcomes are most frequently misclassified. Recall is particularly important for the Dropout category because low recall means that some students who eventually drop out are not identified by the model. Precision indicates how often students predicted to belong to a particular category actually belong to that category. Low precision may result in unnecessary or inappropriate support alerts.

The final report should include the actual confusion matrix, class-wise classification metrics, overall accuracy, and macro-averaged scores produced by the notebook. If a second model uses first-semester academic performance, its results must be interpreted separately because it has access to information unavailable at the time of enrollment.

5. Interpretation and Recommendations

5.1 Interpretation of Model Information

Logistic regression coefficients indicate how predictor variables are associated with the log-odds of a particular student outcome, while other included predictors are held constant.

For one-hot encoded categorical variables, a coefficient represents the difference relative to the reference category. For standardized numerical variables, the coefficient represents the change associated with a one-standard-deviation increase in the predictor.

Positive and negative coefficients indicate the direction of the association with a particular outcome relative to the model's reference class. Their interpretation depends on the coding of the target variable and the predictors included in the model.

The notebook identifies the coefficients with the largest absolute values for each outcome class. These coefficients can help identify variables associated with Dropout, Enrolled, and Graduate predictions. However, coefficient magnitude and predictive importance do not establish causation. The findings should be interpreted in the context of the dataset and the model's limitations.

5.2 Evidence-Based Recommendations

1. **Provide targeted student support:** Use model predictions as an initial indicator for offering voluntary academic advising, tutoring, mentoring, and financial-aid guidance. Predictions should not be used as the sole basis for decisions affecting students.

2. **Monitor class-specific performance:** Regularly review precision, recall, and macro-averaged F1-score. Examine false-negative predictions for Dropout and Enrolled students to understand which students may be overlooked by the model.

3. **Promote fairness and transparency:** Evaluate model performance across relevant student groups, protect personal information, explain how predictions are used, and provide human review and a process for correcting inaccurate records.

4. **Validate the model regularly:** Test the model on data from later student cohorts and reassess its performance after changes to academic programmes, institutional policies, or student-support services. Monitor changes in the data and evaluate whether interventions improve student outcomes.

5. **Match predictions to the intervention stage:** Use information available at enrollment for admission-stage or early-support decisions. Use first-semester academic results only when providing support after that semester. This prevents the use of information that would not have been available at the intended prediction time.

5.3 Limitations

The study has several limitations. First, the dataset represents a particular higher education institution, so the findings may not generalize to other institutions or countries. Second, the outcome classes are imbalanced, which can affect the model's ability to identify students in the smaller categories.

Third, administrative and socioeconomic variables may reflect existing structural disadvantages. Their association with predicted outcomes should not be interpreted as evidence that these characteristics cause students to drop out. Fourth, logistic regression may not capture complex nonlinear relationships or interactions between variables.

Finally, the dataset is observational, so it cannot establish causal relationships. Even if the model achieves good predictive performance, this does not prove that a particular intervention will improve student outcomes. The model should therefore be piloted, evaluated for fairness, and used alongside human judgment rather than as an automatic decision-making system.

6. Conclusion

This case study demonstrates how data science techniques can be used to investigate student academic outcomes and develop a multiclass classification model for predicting Dropout, Enrolled, and Graduate status.

The analysis includes data preprocessing, exploratory data analysis, descriptive statistics, hypothesis testing, logistic regression, and model evaluation. It emphasizes the importance of selecting predictors that are available at the intended prediction time, checking statistical assumptions, and evaluating performance using class-wise and macro-averaged metrics.

The findings can help institutions identify patterns associated with student outcomes and investigate appropriate support initiatives. However, predictions should be treated as indicators for further assessment rather than definitive judgments about individual students. Responsible implementation requires regular validation, fairness monitoring, privacy protection, and human oversight.